

analysis

August 15, 2025

0.1 Hybrid Model Evaluation

Jakob Balkovec Fri, Aug 15th 2025

0.2 Data

```
[2]: import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import pandas as pd

CSV_TRAIN_FILE = "/Users/jbalkovec/Desktop/DR/experiments/model_logs/
↳hybrid_train_history.csv"
CSV_VAL_FILE = "/Users/jbalkovec/Desktop/DR/experiments/model_logs/
↳hybrid_val_history.csv"
```

```
[3]: train_df = pd.read_csv(CSV_TRAIN_FILE)
train_df.head()
```

```
[3]:      loss  f1_micro  roc_auc  subset_acc  hamming  iou_score  dice_score  \
0   1.696087  0.542398      0       0.1     0.400    0.558333  0.689048
1   1.659583  0.633307      0       0.2     0.350    0.600000  0.716667
2   1.635090  0.828151      0       0.6     0.125    0.833333  0.890000
3   1.604202  0.784636      0       0.4     0.150    0.783333  0.866667
4   1.542188  0.797794      0       0.5     0.125    0.833333  0.900000

      pearson_r  epoch
0   -0.160528      1
1    0.534618      2
2    0.693521      3
3    0.699530      4
4    0.739126      5
```

```
[4]: val_df = pd.read_csv(CSV_VAL_FILE)
val_df.head()
```

```
[4]:      loss  f1_micro  roc_auc  subset_acc  hamming  iou_score  dice_score  \
0   1.656705  0.542398      0       0.1     0.400    0.558333  0.689048
1   1.661326  0.570175      0       0.2     0.325    0.625000  0.735714
```

2	1.649466	0.584064	0	0.3	0.250	0.683333	0.782381
3	1.633023	0.584064	0	0.3	0.250	0.683333	0.782381
4	1.605927	0.459064	0	0.3	0.275	0.658333	0.763333

	pearson_r	epoch
0	0.367985	1
1	0.508813	2
2	0.602116	3
3	0.649964	4
4	0.676452	5

```
[20]: _METRIC_ALIASES = {
    "f1": "f1_micro",
    "f1_micro": "f1_micro",
    "acc": "subset_acc",
    "subset_acc": "subset_acc",
    "ham": "hamming",
    "hamming": "hamming",
    "loss": "loss",
    "r": "pearson_r",
    "pearson_r": "pearson_r",
    "dice": "dice_score",
    "dice_score": "dice_score",
    "iou": "iou_score",
    "iou_score": "iou_score",
    "auc": "roc_auc",
    "roc_auc": "roc_auc",
}

_DEFAULT_BANDS = {
    "f1_micro": [
        (0.0, 0.60, "red",      "Poor (< 0.60)"),
        (0.60, 0.75, "yellow", "Moderate (0.60-0.75)"),
        (0.75, 1.00, "green",   "Good (> 0.75)"),
    ],
    "subset_acc": [
        (0.0, 0.40, "red",      "Low (< 0.40)"),
        (0.40, 0.65, "yellow", "Moderate (0.40-0.65)"),
        (0.65, 1.00, "green",   "Good (> 0.65)"),
    ],
    "hamming": [
        (0.00, 0.15, "green",   "Low Error (< 0.15)"),
        (0.15, 0.25, "yellow", "Moderate Error (0.15-0.25)"),
        (0.25, 1.00, "red",     "High Error (> 0.25)"),
    ],
    "pearson_r": [
        (-1.00, 0.30, "red",    "Weak (< 0.30)"),
    ],
}
```

```

        (0.30, 0.60, "yellow", "Moderate (0.30-0.60)"),
        (0.60, 1.00, "green", "Strong (> 0.60)),
    ],
    "dice_score": [
        (0.00, 0.60, "red",      "Poor (< 0.60)"),
        (0.60, 0.75, "yellow", "Fair (0.60-0.75)"),
        (0.75, 1.00, "green",   "Good (> 0.75)"),
    ],
    "iou_score": [
        (0.00, 0.50, "red",      "Low (< 0.50)"),
        (0.50, 0.70, "yellow", "Moderate (0.50-0.70)"),
        (0.70, 1.00, "green",   "High (> 0.70)"),
    ],
    "roc_auc": [
        (0.50, 0.70, "red",      "Below target (< 0.70)"),
        (0.70, 0.85, "yellow", "Decent (0.70-0.85)"),
        (0.85, 1.00, "green",   "Strong (> 0.85)"),
    ],
},
}

_BOUNDED_01 = {"f1_micro", "subset_acc", "hamming", "pearson_r", "dice_score", ↴
    "iou_score", "roc_auc"}
```

`def _norm_metric_name(metric: str) -> str:`

```
key = metric.strip().lower()
if key not in _METRIC_ALIASES:
    raise ValueError(f"Unknown metric '{metric}'. Known:{sorted(_METRIC_ALIASES.keys())}")
return _METRIC_ALIASES[key]
```

`def _loss_bands_for_lambda(loss_lambda: float):`

```
"""
Heuristic bands for total_loss = BCEWithLogitsLoss + * DiceLoss.
Calibrated so that for =1, 'moderate' covers roughly 0.9-1.8 and 'high' >1.8.
"""

# Base BCE bands (no segmentation): low<0.7, moderate 0.7-1.5, high>1.5
# Shift upward by the expected segmentation contribution 0.4* (mid
    ↴DiceLoss)
low_thr = 0.7 + 0.4 * loss_lambda
high_thr = 1.5 + 0.4 * loss_lambda

# Keep some headroom for rougher phases
max_thr = high_thr + 1.0 # only for band shading upper bound
```

`return [`

```
(0.00, low_thr, "green", f"Low (< {low_thr:.2f})"),
```

```

        (low_thr, high_thr, "yellow", f"Moderate ({low_thr:.2f}-{high_thr:.
˓→2f})"),
        (high_thr, max_thr, "red",     f"High (> {high_thr:.2f})"),
    ]

def plot_metric(df_train,
                 df_val,
                 metric: str,
                 *,
                 bands: dict | None = None,
                 show_trend: bool = True,
                 loss_lambda: float = 1.0):

    # desc: plot a metric for train and validation using two DataFrames.

    # arg1: df_train, df_val - DataFrames with columns ['epoch', metric]
    # arg2: metric - name or alias (e.g., 'iou', 'dice', 'f1', 'acc', 'auc', 'r',
    # ˓→'hamming', 'loss')
    # arg3: bands - optional dict to override or extend threshold bands
    # arg4: show_trend - fit a quadratic trend on validation and draw it

    m = _norm_metric_name(metric)
    label = m.replace('_', ' ').title()

    if "epoch" not in df_train or "epoch" not in df_val:
        raise ValueError("Both DataFrames must include an 'epoch' column.")

    if m not in df_train or m not in df_val:
        raise ValueError(f"Metric '{m}' not found in DataFrames.")

    band_cfg = dict(_DEFAULT_BANDS)
    if bands:
        band_cfg.update(bands)

    fig, axes = plt.subplots(2, 1, figsize=(10, 10), sharex=True)

    ax_tr = axes[0]
    ax_tr.plot(df_train["epoch"], df_train[m], marker='o', label='Train',
    # linewidth=2)
    ax_tr.set_ylabel(label, fontsize=12)
    ax_tr.set_title(f"{label} - Train", fontsize=14)
    ax_tr.grid(True, linestyle='--', alpha=0.5)

    ax_va = axes[1]
    ax_va.plot(df_val["epoch"], df_val[m], marker='o', label='Validation',
    # linewidth=2)
    ax_va.set_xlabel("Epoch", fontsize=12)

```

```

ax_va.set_ylabel(label, fontsize=12)
ax_va.set_title(f"{label} - Validation", fontsize=14)
ax_va.grid(True, linestyle='--', alpha=0.5)

if show_trend and df_val[m].notna().sum() >= 3:
    try:
        z = np.polyfit(df_val["epoch"].to_numpy(), df_val[m].to_numpy(), 2)
        p = np.poly1d(z)
        x_vals = np.linspace(df_val["epoch"].min(), df_val["epoch"].max(), ↴
200)
        ax_va.plot(x_vals, p(x_vals), "--", label="Trend")
    except Exception:
        pass # be robust

def draw_bands(ax):

    if m == "loss":
        for lo, hi, color, lab in _loss_bands_for_lambda(loss_lambda):
            ax.axhspan(lo, hi, color=color, alpha=0.10, label=lab)
        return

    if m in band_cfg:
        for lo, hi, color, lab in band_cfg[m]:
            ax.axhspan(lo, hi, color=color, alpha=0.10, label=lab)

draw_bands(ax_tr)
draw_bands(ax_va)

if m in _BOUNDED_01:
    ax_tr.set_ylim(0.0, 1.0)
    ax_va.set_ylim(0.0, 1.0)
elif m == "pearson_r":
    ax_tr.set_ylim(-1.0, 1.0)
    ax_va.set_ylim(-1.0, 1.0)
elif m == "loss":
    # auto with margin
    for ax in (ax_tr, ax_va):
        vals = np.concatenate([df_train[m].dropna().to_numpy(), df_val[m].dropna().to_numpy()])
        if vals.size:
            lo, hi = float(np.nanmin(vals)), float(np.nanmax(vals))
            span = max(1e-6, hi - lo)
            ax.set_ylim(max(0.0, lo - 0.1 * span), hi + 0.1 * span)

val_clean = df_val[df_val[m].notna()]
if not val_clean.empty:
    if m in {"loss", "hamming"}:

```

```

        best_idx = val_clean[m].idxmin()
    else:
        best_idx = val_clean[m].idxmax()
best_epoch = val_clean.loc[best_idx, "epoch"]
best_value = val_clean.loc[best_idx, m]
ax_va.scatter(best_epoch, best_value, s=100, zorder=5, label=f"Best:{best_value:.3f}")
ax_va.axvline(best_epoch, linestyle=":", linewidth=1.5)
ax_va.text(best_epoch, best_value, f" {best_value:.3f}", fontsize=10, va="bottom")

def dedup_legend(ax):
    handles, labels = ax.get_legend_handles_labels()
    seen = []
    new_h, new_l = [], []
    for h, l in zip(handles, labels):
        if l not in seen:
            seen[l] = True
            new_h.append(h)
            new_l.append(l)
    ax.legend(new_h, new_l, loc="best")


dedup_legend(ax_tr)
dedup_legend(ax_va)

plt.tight_layout()
plt.show()

```


0.3 F1

[]: plot_metric(train_df, val_df, "f1")




0.4 Hamming Loss

```
[9]: plot_metric(train_df, val_df, "hamming")
```




0.5 Subset Accuracy

```
[10]: plot_metric(train_df, val_df, "subset_acc")
```




0.6 Loss

```
[21]: plot_metric(train_df, val_df, "loss")
```




0.7 Pearson Correlation (R)

```
[12]: plot_metric(train_df, val_df, "r")
```



0.8 Dice

```
[13]: plot_metric(train_df, val_df, "dice")
```



0.9 IoU

```
[14]: plot_metric(train_df, val_df, "iou")
```

