

analysisV3

August 17, 2025

0.1 Hybrid Model Evaluation

Jakob Balkovec Fri, Aug 15th 2025

0.2 Data

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import pandas as pd

METRICS_LOG_CSV = "metrics_log_latest.csv"
NPZ_DIR = "npz"
MASK_VISUALS_DIR = "mask_visuals"

LOSS_LAMBDA = 1.0
```

```
[2]: df = pd.read_csv(METRICS_LOG_CSV)
df.head()
```

```
[2]:   epoch      lr  train_loss  val_loss  train_f1_micro  train_f1_macro \
0      1  0.000010    0.679842   19.476119      0.898617      0.886766
1      2  0.000010    0.667445   38.415935      0.915791      0.903697
2      3  0.000010    0.517392   52.851267      0.983686      0.979669
3      4  0.000005    0.564369   87.611669      0.997005      0.995839
4      5  0.000005    0.637261  108.560271      0.998608      0.997989

      train_precision_micro  train_recall_micro  train_iou_micro \
0                  0.816809        0.998634      0.815898
1                  0.845540        0.998772      0.844662
2                  0.969635        0.998152      0.967897
3                  0.995312        0.998703      0.994027
4                  0.997714        0.999503      0.997220

      train_roc_auc_macro ...  train_coord_r2  val_coord_mae  val_coord_rmse \
0             NaN ...       -4.600443      0.325995      0.400923
1             NaN ...       -4.840408      0.315038      0.386400
2             NaN ...       -5.996484      0.283642      0.350308
```

```

3          NaN ...      -6.884182      0.300091      0.366335
4          NaN ...      -7.215684      0.301701      0.368077

  val_coord_r2  train_seg_iou  train_seg_dice  val_seg_iou  val_seg_dice \
0    -1.310547      0.021654      0.042390      0.003729      0.007431
1    -1.146186      0.025959      0.050604      0.004036      0.008040
2    -0.763979      0.026841      0.052279      0.004305      0.008572
3    -0.929079      0.028381      0.055196      0.004210      0.008385
4    -0.947465      0.029913      0.058089      0.003814      0.007599

      monitor  monitor_value
0  val_f1_micro      0.756196
1  val_f1_micro      0.683835
2  val_f1_micro      0.681402
3  val_f1_micro      0.681213
4  val_f1_micro      0.681213

[5 rows x 60 columns]

```

```
[3]: def split_train_val(df):
    base = pd.DataFrame({"epoch": df["epoch"]})

    train_cols = [c for c in df.columns if c.startswith("train_")]
    val_cols   = [c for c in df.columns if c.startswith("val_")]

    def strip_prefix(cols, prefix):
        out = {}
        for c in cols:
            key = c[len(prefix):]
            out[key] = c
        return out

    train_map = strip_prefix(train_cols, "train_")
    val_map   = strip_prefix(val_cols, "val_")

    shared = sorted(set(train_map.keys()) & set(val_map.keys()))

    df_train = base.copy()
    df_val   = base.copy()
    for k in shared:
        df_train[k] = df[train_map[k]]
        df_val[k]   = df[val_map[k]]

    return df_train, df_val, shared

df_train, df_val, shared_metrics = split_train_val(df)
```

```
[4]: df_train.head()
```

```
[4]:    epoch  coord_mae  coord_r2  coord_rmse  f1_class_0  f1_class_1  f1_class_2  \
0       1   0.510571 -4.600443   0.622578   0.874932   0.994936   0.965306
1       2   0.519847 -4.840408   0.635766   0.885020   0.996473   0.971634
2       3   0.569415 -5.996484   0.695858   0.975304   0.999136   0.994531
3       4   0.600906 -6.884182   0.738680   0.997633   0.999933   0.999546
4       5   0.613254 -7.215684   0.754053   0.999339   0.999978   0.999916

      f1_class_3  f1_macro  f1_micro  ...  precision_class_3  precision_micro  \
0   0.711892   0.886766   0.898617  ...   0.553359   0.816809
1   0.761660   0.903697   0.915791  ...   0.617613   0.845540
2   0.949704   0.979669   0.983686  ...   0.912570   0.969635
3   0.986244   0.995839   0.997005  ...   0.979902   0.995312
4   0.992722   0.997989   0.998608  ...   0.988414   0.997714

  recall_class_0  recall_class_1  recall_class_2  recall_class_3  \
0   0.999166   0.997686   0.999713   0.997738
1   0.999626   1.000000   0.999857   0.993339
2   0.999195   1.000000   0.999976   0.989987
3   0.999684   1.000000   0.999952   0.992669
4   0.999942   1.000000   1.000000   0.997068

  recall_micro  roc_auc_macro  seg_dice  seg_iou
0   0.998634           NaN  0.042390  0.021654
1   0.998772           NaN  0.050604  0.025959
2   0.998152           NaN  0.052279  0.026841
3   0.998703           NaN  0.055196  0.028381
4   0.999503           NaN  0.058089  0.029913
```

[5 rows x 29 columns]

```
[5]: df_val.head()
```

```
[5]:    epoch  coord_mae  coord_r2  coord_rmse  f1_class_0  f1_class_1  f1_class_2  \
0       1   0.325995 -1.310547   0.400923   0.876564   0.867433   0.600865
1       2   0.315038 -1.146186   0.386400   0.876564   0.867433   0.027227
2       3   0.283642 -0.763979   0.350308   0.876564   0.867433   0.001957
3       4   0.300091 -0.929079   0.366335   0.876564   0.867433   0.000000
4       5   0.301701 -0.947465   0.368077   0.876564   0.867433   0.000000

      f1_class_3  f1_macro  f1_micro  ...  precision_class_3  precision_micro  \
0       0.0   0.586215   0.756196  ...       0.0       0.782517
1       0.0   0.442806   0.683835  ...       0.0       0.774027
2       0.0   0.436488   0.681402  ...       0.0       0.773150
3       0.0   0.435999   0.681213  ...       0.0       0.773076
4       0.0   0.435999   0.681213  ...       0.0       0.773076
```

```

recall_class_0  recall_class_1  recall_class_2  recall_class_3 \
0              1.0            1.0          0.469916      0.0
1              1.0            1.0          0.013806      0.0
2              1.0            1.0          0.000979      0.0
3              1.0            1.0          0.000000      0.0
4              1.0            1.0          0.000000      0.0

recall_micro   roc_auc_macro  seg_dice    seg_iou
0      0.731588           NaN  0.007431  0.003729
1      0.612469           NaN  0.008040  0.004036
2      0.609119           NaN  0.008572  0.004305
3      0.608863           NaN  0.008385  0.004210
4      0.608863           NaN  0.007599  0.003814

```

[5 rows x 29 columns]

```

[ ]: def LOSS_BANDS(loss_lambda: float = 1.0):
    low_thr  = 0.7 + 0.4 * loss_lambda
    high_thr = 1.5 + 0.4 * loss_lambda
    max_thr  = high_thr + 1.0
    return [
        (0.00,  low_thr, "green", f"Low (< {low_thr:.2f})"),
        (low_thr, high_thr,"yellow", f"Moderate ({low_thr:.2f}-{high_thr:.2f})"),
        (high_thr, max_thr,"red",   f"High (> {high_thr:.2f})"),
    ]

BANDS = {
    "f1_micro": [
        (0.00, 0.60, "red",      "Poor (< 0.60)"),
        (0.60, 0.75, "yellow",   "Moderate (0.60-0.75)"),
        (0.75, 1.00, "green",    "Good (> 0.75)"),
    ],
    "f1_macro": [
        (0.00, 0.60, "red",      "Poor (< 0.60)"),
        (0.60, 0.75, "yellow",   "Moderate (0.60-0.75)"),
        (0.75, 1.00, "green",    "Good (> 0.75)"),
    ],
    "precision_micro": [
        (0.00, 0.60, "red",      "Low (< 0.60)"),
        (0.60, 0.80, "yellow",   "Moderate (0.60-0.80)"),
        (0.80, 1.00, "green",    "High (> 0.80)"),
    ],
    "recall_micro": [
        (0.00, 0.60, "red",      "Low (< 0.60)"),
        (0.60, 0.80, "yellow",   "Moderate (0.60-0.80)"),

```

```

        (0.80, 1.00, "green", "High (> 0.80")),
    ],
    "iou_micro": [
        (0.00, 0.50, "red", "Low (< 0.50)"),
        (0.50, 0.70, "yellow", "Moderate (0.50-0.70)"),
        (0.70, 1.00, "green", "High (> 0.70)"),
    ],
    "roc_auc_macro": [
        (0.50, 0.70, "red", "Below target (< 0.70)"),
        (0.70, 0.85, "yellow", "Decent (0.70-0.85)"),
        (0.85, 1.00, "green", "Strong (> 0.85)"),
    ],
    "seg_iou": [
        (0.00, 0.50, "red", "Low (< 0.50)"),
        (0.50, 0.70, "yellow", "Moderate (0.50-0.70)"),
        (0.70, 1.00, "green", "High (> 0.70)"),
    ],
    "seg_dice": [
        (0.00, 0.60, "red", "Poor (< 0.60)"),
        (0.60, 0.75, "yellow", "Fair (0.60-0.75)"),
        (0.75, 1.00, "green", "Good (> 0.75)"),
    ],
    "coord_r2": [
        (-5.00, 0.00, "red", "Worse than baseline (< 0)"),
        (0.00, 0.50, "yellow", "Moderate (0-0.5)"),
        (0.50, 1.00, "green", "Good (> 0.5)"),
    ],
    "coord_mae": [
        (0.20, 10.0, "red", "High error (> 0.20)"),
        (0.10, 0.20, "yellow", "Moderate (0.10-0.20)"),
        (0.00, 0.10, "green", "Low (< 0.10)"),
    ],
    "coord_rmse": [
        (0.25, 10.0, "red", "High error (> 0.25)"),
        (0.15, 0.25, "yellow", "Moderate (0.15-0.25)"),
        (-5.0, 0.15, "green", "Low (< 0.15)"),
    ],
},
}

for c in range(4):
    BANDS[f"f1_class_{c}"] = [
        (0.00, 0.60, "red", "Poor (< 0.60)"),
        (0.60, 0.75, "yellow", "Moderate (0.60-0.75)"),
        (0.75, 1.00, "green", "Good (> 0.75)"),
    ]

```

```

BANDS[f"precision_class_{c}"] = [
    (0.00, 0.60, "red", "Low (< 0.60)"),
    (0.60, 0.80, "yellow", "Moderate (0.60-0.80)"),
    (0.80, 1.00, "green", "High (> 0.80)"),
]
BANDS[f"recall_class_{c}"] = [
    (0.00, 0.60, "red", "Low (< 0.60)"),
    (0.60, 0.80, "yellow", "Moderate (0.60-0.80)"),
    (0.80, 1.00, "green", "High (> 0.80)"),
]
BANDS[f"iou_class_{c}"] = [
    (0.00, 0.50, "red", "Low (< 0.50)"),
    (0.50, 0.70, "yellow", "Moderate (0.50-0.70)"),
    (0.70, 1.00, "green", "High (> 0.70)"),
]

def get_bands(metric_name: str, *, loss_lambda: float = 1.0):
    if metric_name == "loss":
        return LOSS_BANDS(loss_lambda)
    return BANDS.get(metric_name, None)

```

```

[ ]: def plot_train_val(df, metric_name, show_trend=True):
    train_col = f"train_{metric_name}"
    val_col = f"val_{metric_name}"
    if train_col not in df.columns or val_col not in df.columns:
        raise ValueError(f"{metric_name} not found in DataFrame.")

    x = df["epoch"]
    train_y = df[train_col]
    val_y = df[val_col]

    fig, axes = plt.subplots(2, 1, figsize=(10, 8), sharex=True)

    axes[0].plot(x, train_y, marker='o', label="Train", linewidth=2)
    axes[0].set_ylabel(metric_name)
    axes[0].set_title(f"{metric_name} - Train")
    axes[0].grid(True, linestyle="--", alpha=0.5)

    axes[1].plot(x, val_y, marker='o', label="Validation", linewidth=2)
    axes[1].set_xlabel("Epoch")
    axes[1].set_ylabel(metric_name)
    axes[1].set_title(f"{metric_name} - Validation")
    axes[1].grid(True, linestyle="--", alpha=0.5)

    if show_trend and val_y.notna().sum() >= 3:
        z = np.polyfit(x, val_y, 2)
        p = np.poly1d(z)

```

```

x_fit = np.linspace(x.min(), x.max(), 200)
axes[1].plot(x_fit, p(x_fit), "--", label="Trend")

bands = get_bands(metric_name, loss_lambda=1.0)
for ax in axes:
    if bands:
        for lo, hi, color, lab in bands:
            ax.axhspan(lo, hi, color=color, alpha=0.10, label=lab)

val_clean = val_y.dropna()
if not val_clean.empty:
    if metric_name in ["loss"]:
        best_idx = val_clean.idxmin()
    else:
        best_idx = val_clean.idxmax()
    best_epoch = x[best_idx]
    best_value = val_clean.loc[best_idx]
    axes[1].scatter(best_epoch, best_value, s=100, zorder=5, label=f"Best:{best_value:.3f}")
    axes[1].axvline(best_epoch, linestyle=":", linewidth=1.5)

for ax in axes:
    handles, labels = ax.get_legend_handles_labels()
    by_label = dict(zip(labels, handles))
    ax.legend(by_label.values(), by_label.keys(), loc="best")


plt.tight_layout()
plt.show()

```

1 Global Classification Metrics


1.1 Loss

[8]: plot_train_val(df, "loss") # fix scale




1.2 F1 - Micro

```
[9]: plot_train_val(df, "f1_micro")
```




1.3 F1 - Macro

```
[10]: plot_train_val(df, "f1_macro")
```




1.4 Percision Micro

```
[11]: plot_train_val(df, "precision_micro")
```




1.5 Recall Micro

```
[12]: plot_train_val(df, "recall_micro")
```



1.6 IoU Micro


```
[13]: plot_train_val(df, "iou_micro")
```




2 Per-Class Classification Metrics

2.1 Micro Aneurysms (MA)


```
[14]: plot_train_val(df, "f1_class_0")
```




```
[15]: plot_train_val(df, "precision_class_0")
```



```
[16]: plot_train_val(df, "recall_class_0")
```




```
[17]: plot_train_val(df, "iou_class_0")
```




2.2 Hemorrhages (HE)


```
[18]: plot_train_val(df, "f1_class_1")
```




```
[19]: plot_train_val(df, "precision_class_1")
```



```
[20]: plot_train_val(df, "recall_class_1")
```




```
[21]: plot_train_val(df, "iou_class_1")
```




2.3 Hard Exudates (EX)


```
[22]: plot_train_val(df, "f1_class_2")
```




```
[23]: plot_train_val(df, "precision_class_2")
```



```
[24]: plot_train_val(df, "recall_class_2")
```




```
[25]: plot_train_val(df, "iou_class_3")
```




2.4 Soft Exudates (SE)


```
[26]: plot_train_val(df, "f1_class_3")
```




```
[27]: plot_train_val(df, "precision_class_3")
```



```
[28]: plot_train_val(df, "recall_class_3")
```




```
[29]: plot_train_val(df, "iou_class_3")
```



2.5 Segmentation Metrics


2.5.1 IoU

```
[30]: plot_train_val(df, "seg_iou")
```



2.6 Dice


```
[31]: plot_train_val(df, "seg_dice")
```



3 Coordinate Regression Metrics


4 MAE

```
[32]: plot_train_val(df, "coord_mae")
```



4.1 RMSE

```
[33]: plot_train_val(df, "coord_rmse")
```



4.2 R^2

```
[34]: plot_train_val(df, "coord_r2")
```

