

Hybrid Model Performance Analysis

August 7, 2025

0.1 Hybrid Model Evaluation

Jakob Balkovec Tue, Aug 5th 2025

This document evaluates the performance of the hybrid model, solely on metrics/numbers I got whilst training it. The model was trained on a dataset of 1841 images, and 184100 patches. Due to the stack mismatch issue, I didn't get a chance to filter out the noise (black patches) but I think the results are still quite promising.

The model is a dual-branch DL architecture (global and local), it leverages two ResNet18 backbones: one processes the full fundus image to extract global context, the other encodes localized lesion patches.

The output features from each branch are then projected into a shared latent space via fully connected layers and fused using a multi-head self-attention mechanism...allowing the model to align and weigh patch-level signals relative to the image context.

The fused representation is passed through a lightweight MLP classifier with a sigmoid activation to predict the presence of four lesion types.

0.2 Training

The model was trained with a learning rate of 0.0001, and I used the Adam optimizer (work to be done here). The training loss decreased steadily over time, and the validation loss also showed a similar trend.

Since I couldn't access Texas A&M's GPU cluster, I had to train the model on Google Colab. I got booted off after about 9.5 epochs, but f1 was still improving, meaning there is performance left on the table.

I designed a new training function that allows for early stopping as well as continuing training from a checkpoint. This way, I can resume training on the cluster without losing progress. I haven't tested any of this yet, but I'm hoping it works...

Google Colab for Education is not available in my country. I emailed Google support, but they said they can't help me. I'll make the most of what I have...

0.3 Data

```
[30]: import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

```
import pandas as pd

CSV_FILE = "/Users/jbalkovec/Desktop/DR/experiments/model_logs/hybrid_model_log.
˓→CSV"
```

```
[31]: df = pd.read_csv(CSV_FILE)
df.head(n=len(df))
```

```
[31]:          timestamp epoch f1_micro hamming subset_acc iou \
0 2025-08-07 06:17:55 1 0.737753 0.241857 0.325733 0.683768
1 2025-08-07 06:25:24 1 0.701047 0.238192 0.315418 0.689558
2 2025-08-07 06:35:14 2 0.754473 0.230185 0.345820 0.690147
3 2025-08-07 06:42:39 2 0.800591 0.211726 0.386536 0.703131
4 2025-08-07 06:52:31 3 0.775229 0.212948 0.387622 0.707338
5 2025-08-07 07:00:05 3 0.813652 0.197476 0.409338 0.727832
6 2025-08-07 07:10:09 4 0.802093 0.200190 0.394680 0.721544
7 2025-08-07 07:17:46 4 0.836930 0.177796 0.461455 0.738916
8 2025-08-07 07:27:44 5 0.818863 0.184853 0.440282 0.734980
9 2025-08-07 07:35:22 5 0.818492 0.167210 0.482628 0.757555
10 2025-08-07 07:45:18 6 0.827449 0.182546 0.454940 0.741450

      dice pearson_r loss
0 0.765708 0.507020 0.509871
1 0.773851 0.556022    NaN
2 0.769050 0.557325 0.474096
3 0.775575 0.607541    NaN
4 0.780621 0.595436 0.444653
5 0.797764 0.630317    NaN
6 0.792883 0.629872 0.421925
7 0.800284 0.671015    NaN
8 0.800610 0.658281 0.399331
9 0.817199 0.711187    NaN
10 0.804788 0.670248 0.388381
```

```
[32]: def plot_metric(df, metric: str):

    # Split into train/val
    df_train = df.iloc[:2].copy()
    df_val = df.iloc[1::2].copy()

    fig, axes = plt.subplots(2, 1, figsize=(10, 10), sharex=True)

    # Plot TRAIN
    ax_train = axes[0]
    ax_train.plot(df_train["epoch"], df_train[metric], marker='o', color='tab:blue', label='Train', linewidth=2)
    ax_train.set_ylabel(metric.replace('_', ' ').title(), fontsize=12)
```

```

    ax_train.set_title(f"{metric.replace('_', ' ')}.title()} - Train", u
    ↪fontsize=14)
    ax_train.grid(True, linestyle='--', alpha=0.5)

    # Plot VAL
    ax_val = axes[1]
    ax_val.plot(df_val["epoch"], df_val[metric], marker='o', color='tab:
    ↪orange', label='Validation', linewidth=2)
    ax_val.set_xlabel("Epoch", fontsize=12)
    ax_val.set_ylabel(metric.replace('_', ' ').title(), fontsize=12)
    ax_val.set_title(f"{metric.replace('_', ' ')}.title()} - Validation", u
    ↪fontsize=14)
    ax_val.grid(True, linestyle='--', alpha=0.5)

    # Trend line on VAL only
    if df_val[metric].notna().sum() >= 3:
        z = np.polyfit(df_val["epoch"], df_val[metric], 2)
        p = np.poly1d(z)
        x_vals = np.linspace(df_val["epoch"].min(), df_val["epoch"].max(), 100)
        ax_val.plot(x_vals, p(x_vals), "--", color="gray", label="Trend Line")

    # ----- THRESHOLDS -----
    def draw_thresholds(ax):
        if metric == 'f1_micro':
            ax.axhspan(0.0, 0.6, color='red', alpha=0.1, label="Poor (< 0.6)")
            ax.axhspan(0.6, 0.75, color='yellow', alpha=0.1, label="Moderate (0.
            ↪6-0.75)")
            ax.axhspan(0.75, 1.0, color='green', alpha=0.1, label="Good (> 0.
            ↪.75)")
        elif metric == 'subset_acc':
            ax.axhspan(0.0, 0.4, color='red', alpha=0.1, label="Low (< 0.4)")
            ax.axhspan(0.4, 0.65, color='yellow', alpha=0.1, label="Moderate (0.
            ↪.4-0.65)")
            ax.axhspan(0.65, 1.0, color='green', alpha=0.1, label="Good (> 0.
            ↪.65)")
        elif metric == 'hamming':
            ax.axhspan(0.0, 0.15, color='green', alpha=0.1, label="Low Error (<
            ↪0.15)")
            ax.axhspan(0.15, 0.25, color='yellow', alpha=0.1, label="Moderate_U
            ↪Error (0.15-0.25)")
            ax.axhspan(0.25, 1.0, color='red', alpha=0.1, label="High Error (>
            ↪0.25)")
        elif metric == 'loss':
            ax.axhspan(0.0, 0.3, color='green', alpha=0.1, label="Low Loss (< 0.
            ↪.3)")


```

```

        ax.axhspan(0.3, 0.6, color='yellow', alpha=0.1, label="Moderate Loss (0.3-0.6)")
        ↵       ax.axhspan(0.6, 1.0, color='red', alpha=0.1, label="High Loss (> 0.6)")

    elif metric == 'pearson_r':
        ax.axhspan(-1.0, 0.3, color='red', alpha=0.1, label="Weak (< 0.3)")
        ax.axhspan(0.3, 0.6, color='yellow', alpha=0.1, label="Moderate (0.3-0.6)")
        ax.axhspan(0.6, 1.0, color='green', alpha=0.1, label="Strong (> 0.6)")

    elif metric == 'dice':
        ax.axhspan(0.0, 0.6, color='red', alpha=0.1, label="Poor (< 0.6)")
        ax.axhspan(0.6, 0.75, color='yellow', alpha=0.1, label="Fair (0.6-0.75)")
        ax.axhspan(0.75, 1.0, color='green', alpha=0.1, label="Good (> 0.75)")

    elif metric == 'iou':
        ax.axhspan(0.0, 0.5, color='red', alpha=0.1, label="Low (< 0.5)")
        ax.axhspan(0.5, 0.7, color='yellow', alpha=0.1, label="Moderate (0.5-0.7)")
        ax.axhspan(0.7, 1.0, color='green', alpha=0.1, label="High (> 0.7)")

draw_thresholds(ax_train)
draw_thresholds(ax_val)

# ----- BEST -----
df_val_clean = df_val[df_val[metric].notna()]
if not df_val_clean.empty:
    if metric in ['loss', 'hamming']:
        best_idx = df_val_clean[metric].idxmin()
    else:
        best_idx = df_val_clean[metric].idxmax()

    best_epoch = df_val_clean.loc[best_idx, "epoch"]
    best_value = df_val_clean.loc[best_idx, metric]
    ax_val.scatter(best_epoch, best_value, color='purple', s=100, zorder=5, label=f"Best: {best_value:.3f}")
    ax_val.text(best_epoch + 0.1, best_value, f"{best_value:.3f}", color='purple', fontsize=10)

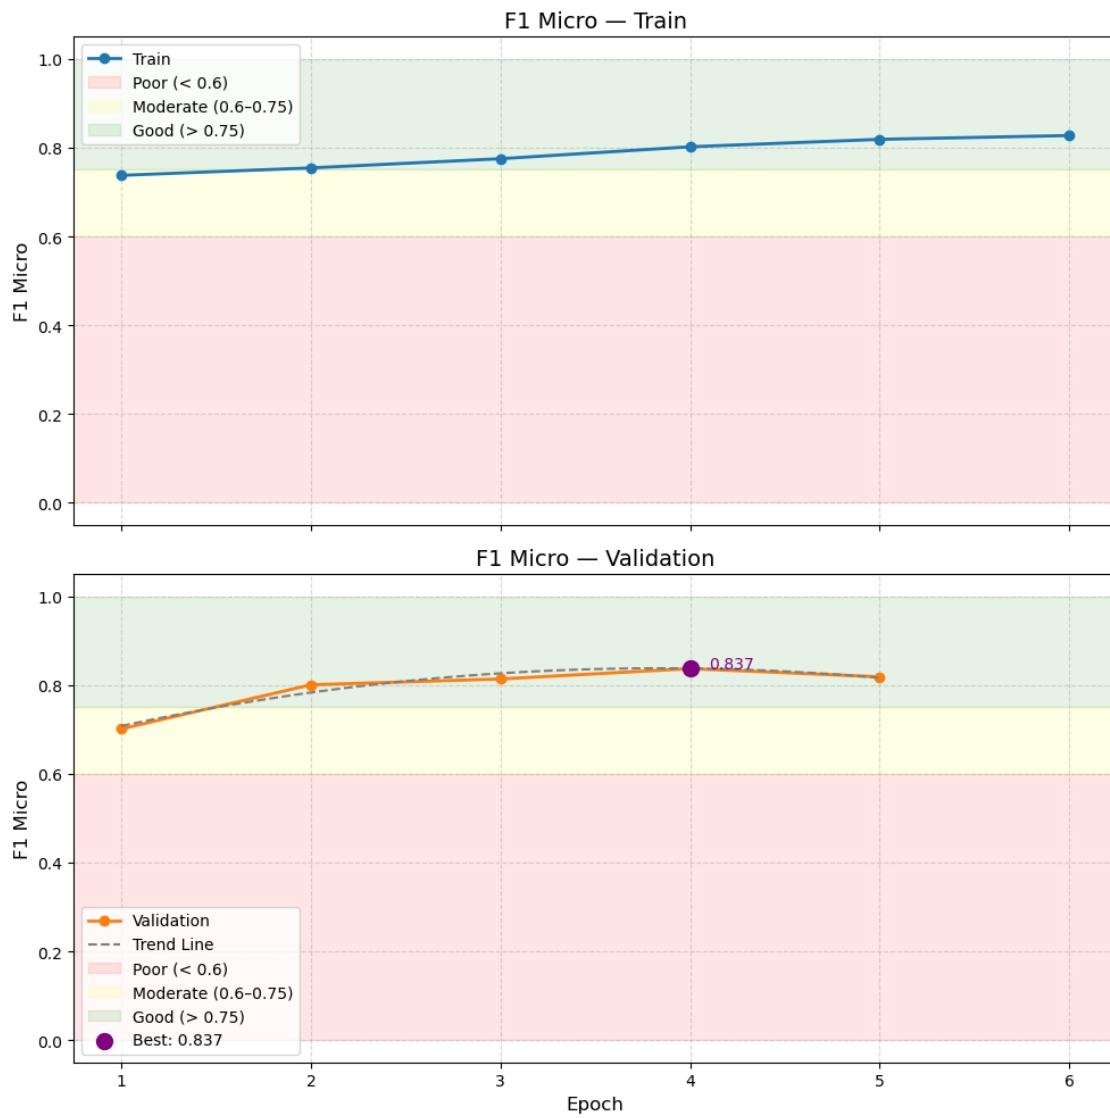
# Legends
ax_train.legend()
ax_val.legend()

plt.tight_layout()
plt.show()

```

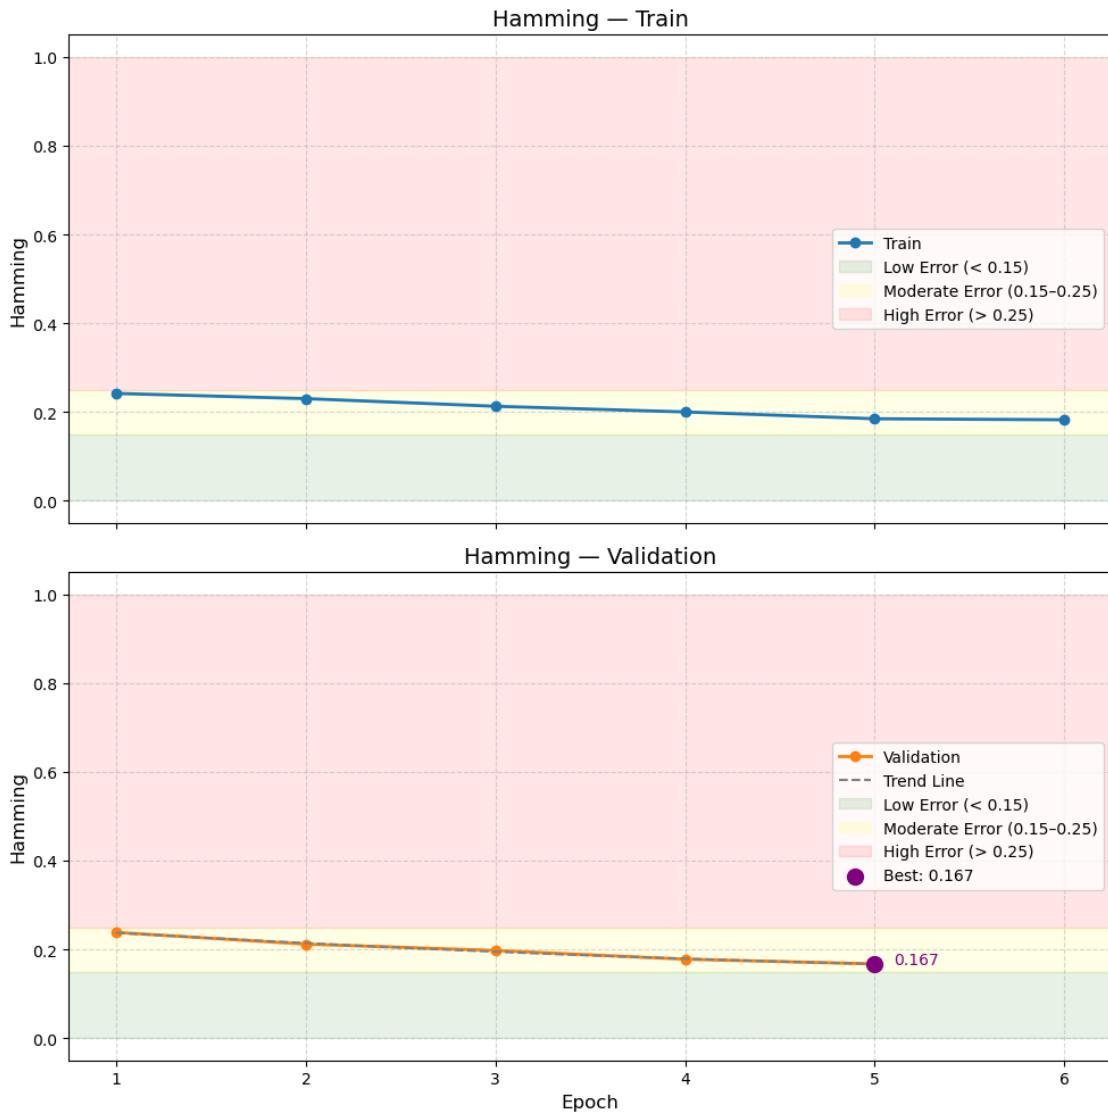
0.4 F1

```
[33]: plot_metric(df, "f1_micro")
```



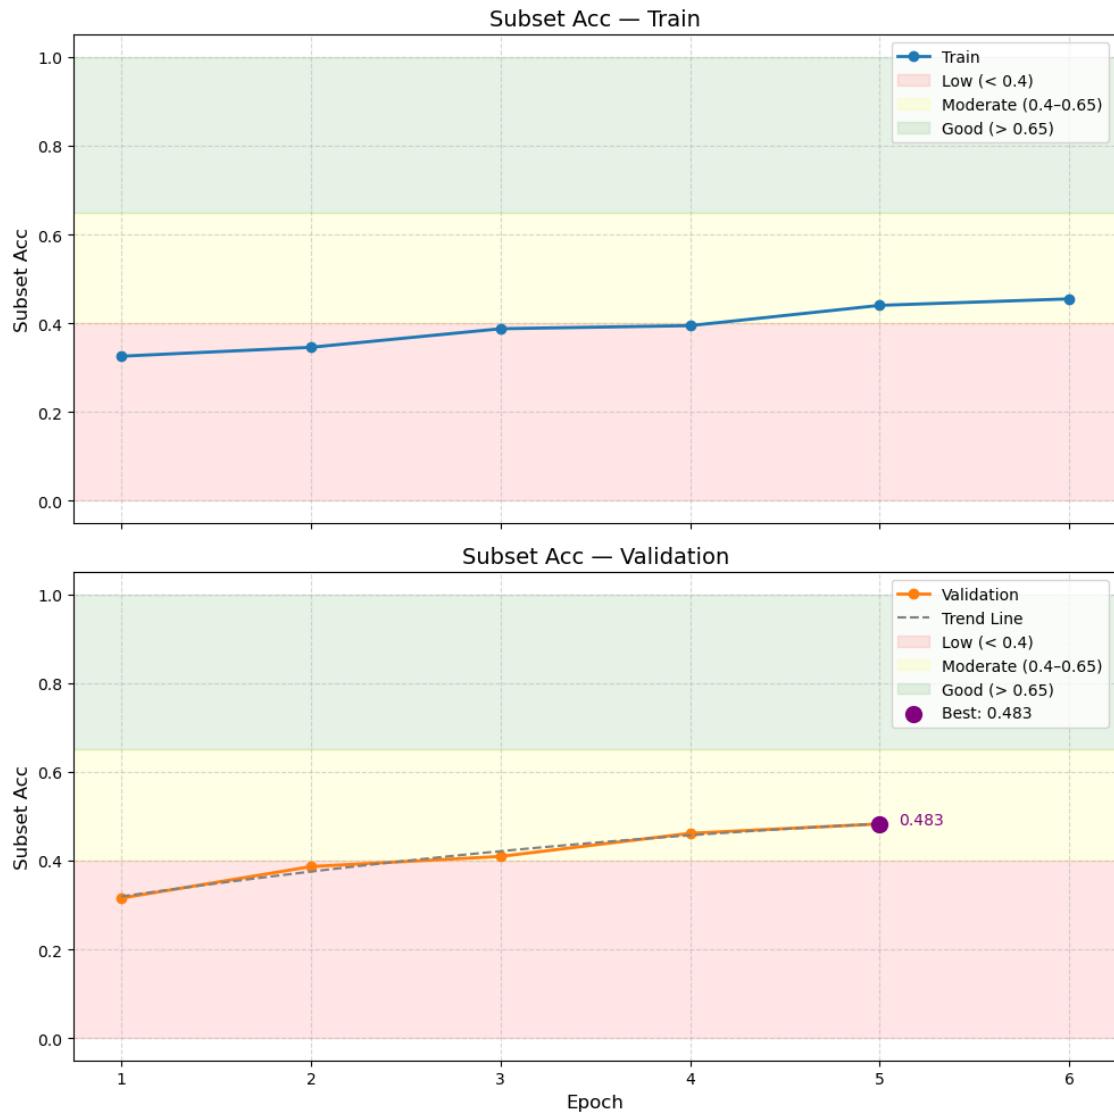
0.5 Hamming Loss

```
[34]: plot_metric(df, "hamming")
```



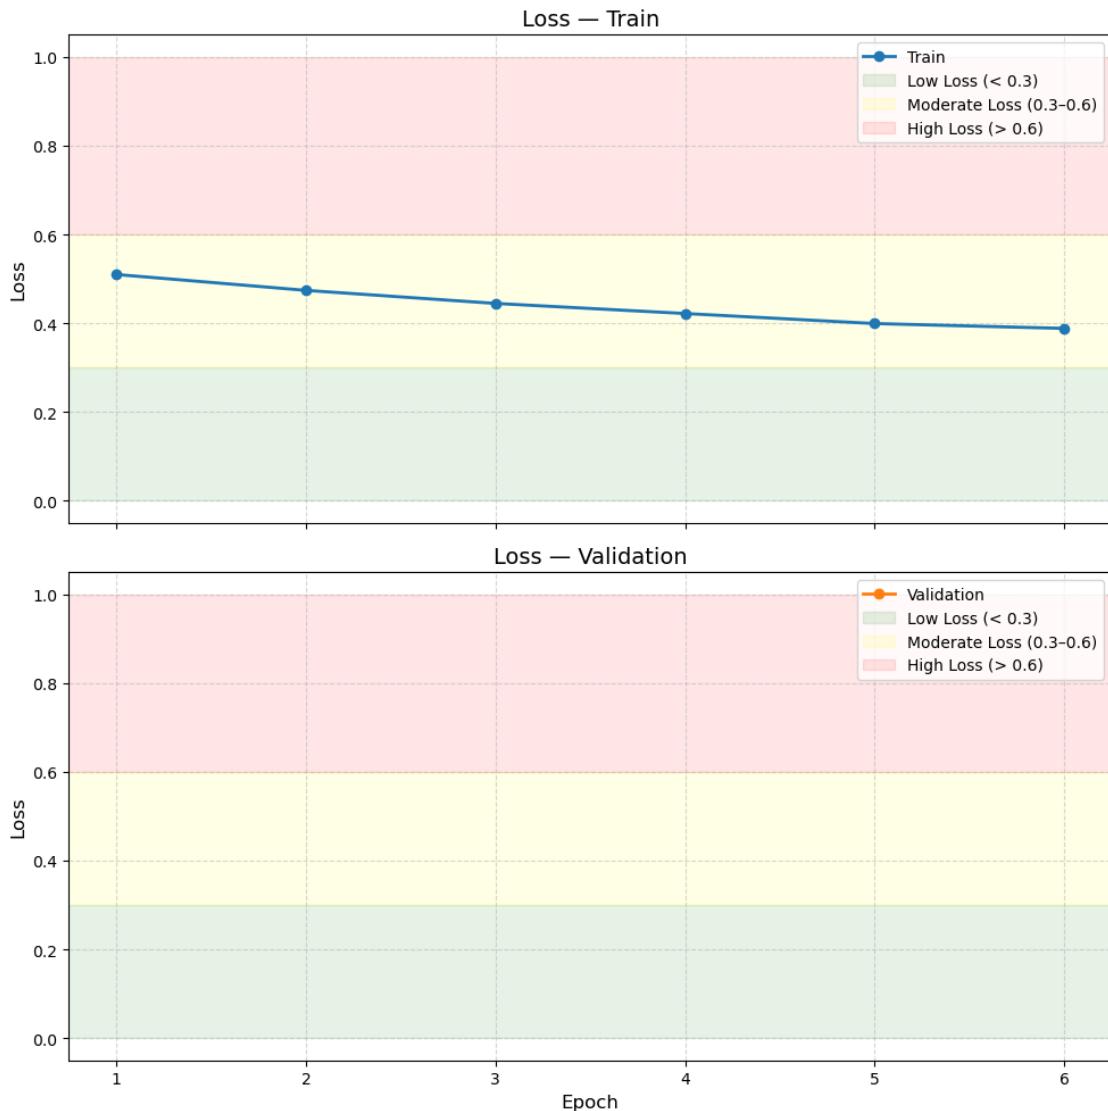
0.6 Subset Accuracy

```
[35]: plot_metric(df, "subset_acc")
```



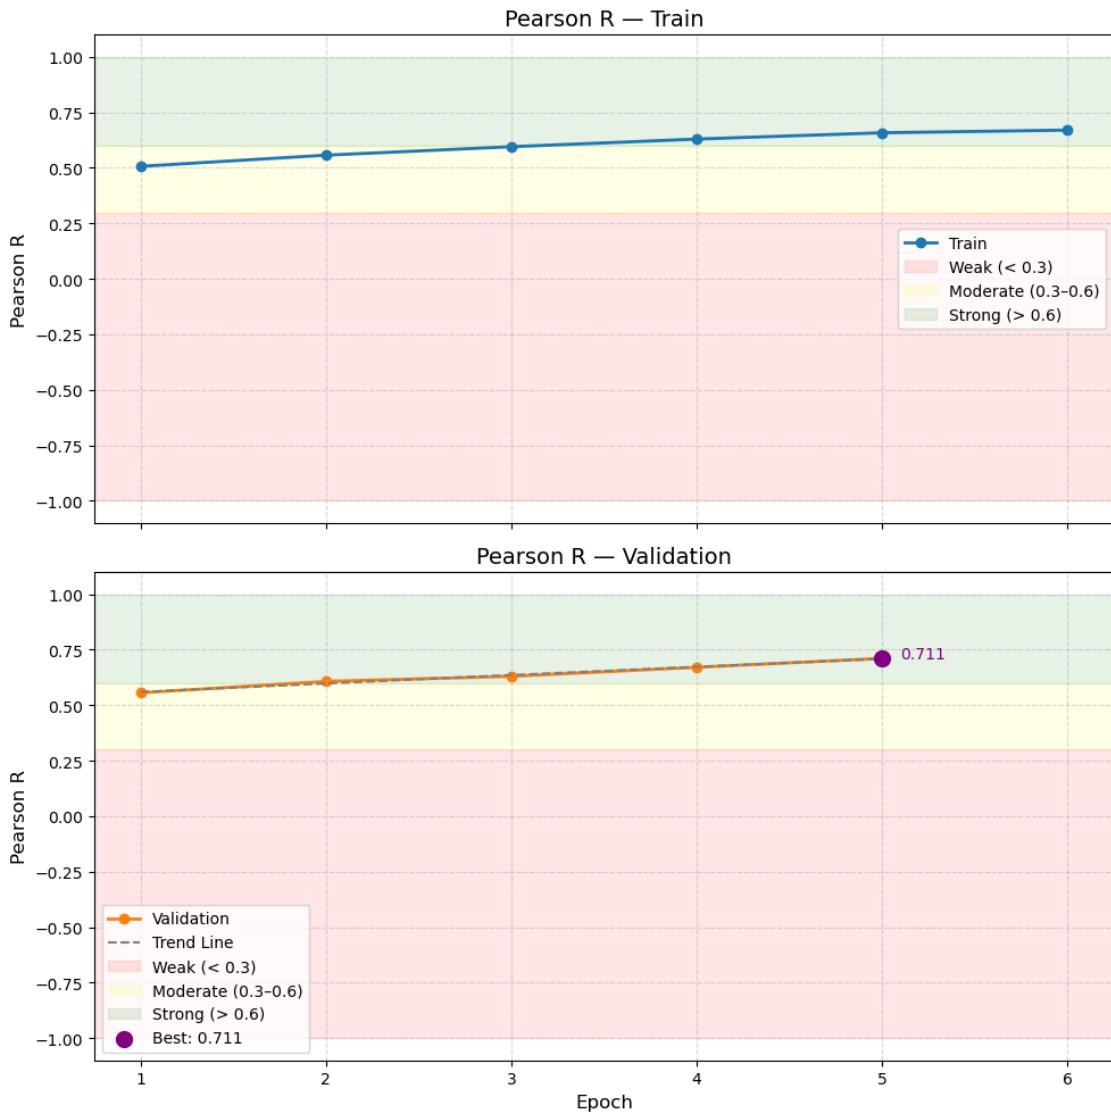
0.7 Loss

```
[36]: plot_metric(df, "loss")
```



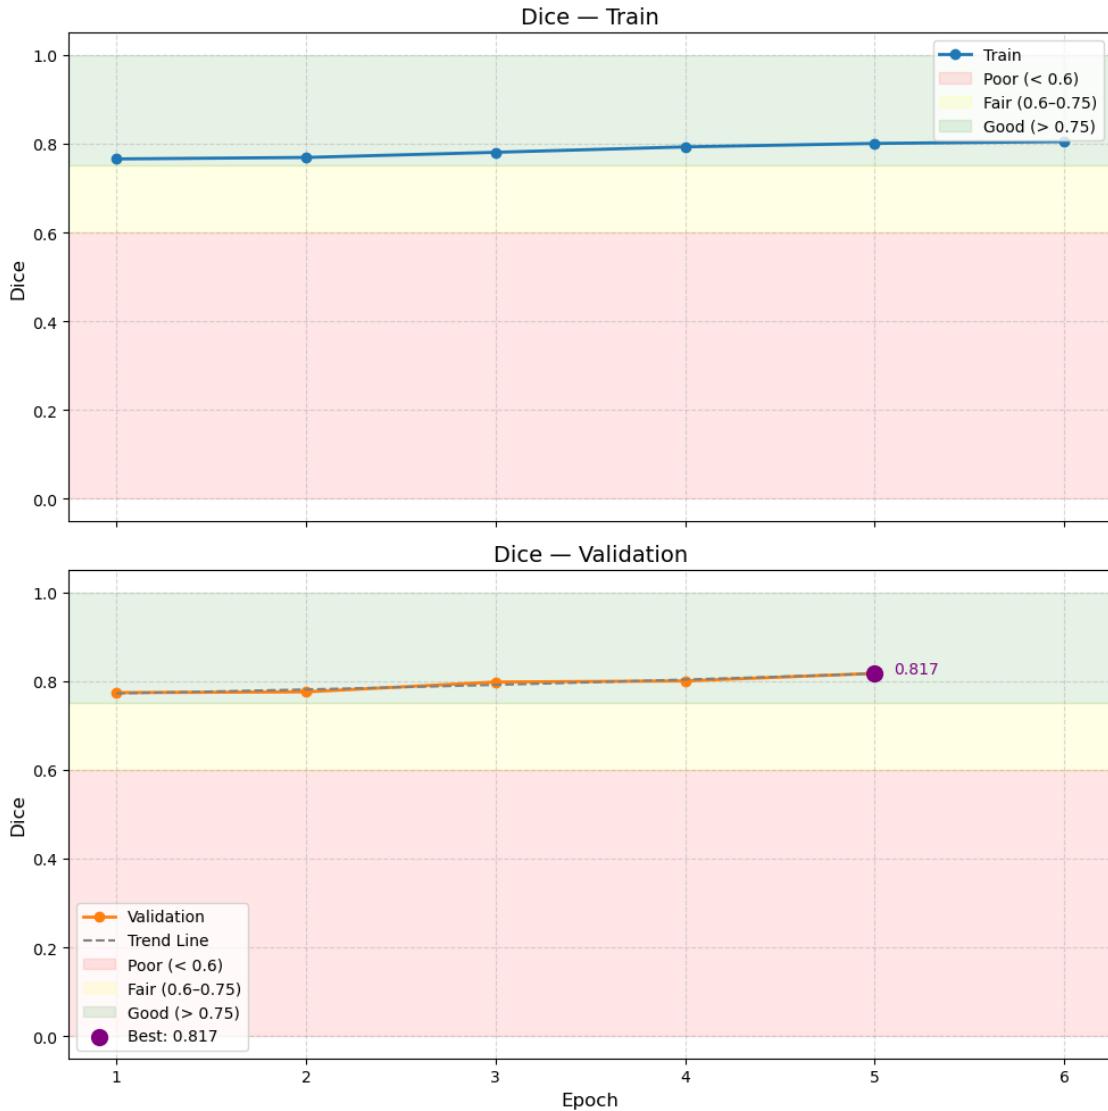
0.8 Pearson Correlation (R)

```
[37]: plot_metric(df, "pearson_r")
```



0.9 Dice

```
[38]: plot_metric(df, "dice")
```



0.10 IoU

```
[39]: plot_metric(df, "iou")
```

